



**Xi'an Jiaotong-Liverpool University**

**西交利物浦大学**

**INT305 Machine Learning  
Lecture 3**

**Linear Classifiers, Logistic Regression, Multiclass Classification**

**Sichen Liu**

**Department Intelligence Science**

**[Sichen.Liu@xjtu.edu.cn](mailto:Sichen.Liu@xjtu.edu.cn)**

## Overview

---

- **Classification**: predicting a discrete-valued target
  - **Binary classification**: predicting a binary-valued target
  - **Multiclass classification**: predicting a discrete( $>2$ )-valued target
- Examples of binary classification
  - Predict whether a patient has a disease, given the presence or absence of various symptoms
  - Classify e-mails as spam or non-spam
  - Predict whether a financial transaction is fraudulent

# Overview

---

## Binary linear classification

- **Classification:** given a  $D$ -dimensional input  $\mathbf{x} \in \mathbb{R}^D$  predict a discrete-valued target
- **Binary:** predict a binary target  $t \in \{0,1\}$ 
  - Training examples with  $t = 1$  are called **positive examples**, and training examples with  $t = 0$  are called **negative examples**. Sorry.
  - $t \in \{0,1\}$  or  $t \in \{-1, +1\}$  is for computational convenience.
- **Linear:** model prediction  $y$  is a linear function of  $\mathbf{x}$ , followed by a threshold  $r$ :

$$z = \mathbf{w}^\top \mathbf{x} + b$$

$$y = \begin{cases} 1 & \text{if } z \geq r \\ 0 & \text{if } z < r \end{cases}$$

## Simplifications

---

### Eliminating the threshold

- We can assume without loss of generality (WLOG) that the threshold  $r = 0$ :

$$\mathbf{w}^\top \mathbf{x} + b \geq r \iff \mathbf{w}^\top \mathbf{x} + b - r \geq 0$$

$\underbrace{\phantom{w^\top x + b - r}}_{\triangleq w_0}$

### Eliminating the bias

- Add a dummy feature  $x_0$  which always takes the value 1. the weight  $w_0 = b$  is equivalent to a bias (same as linear regression)

### Simplified model

- Receive input  $\mathbf{x} \in \mathbb{R}^{D+1}$  with  $x_0 = 1$ :

$$z = \mathbf{w}^\top \mathbf{x}$$

$$y = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

## Examples

---

- Let's consider some simple examples to examine the properties of our model
- Let's focus on minimizing the training set error, and forget about whether our model will generalize to a test set.

## Examples

---

**NOT**

$x_0$	$x_1$	$t$
1	0	1
1	1	0

- Suppose this is our training set, with the dummy feature  $x_0$  included.
- Which conditions on  $w_0, w_1$  guarantee perfect classification?
  - When  $x_1 = 0$ , need:  $z = w_0x_0 + w_1x_1 \geq 0 \iff w_0 \geq 0$
  - When  $x_1 = 1$ , need:  $z = w_0x_0 + w_1x_1 < 0 \iff w_0 + w_1 < 0$
- Example solution:  $w_0 = 1, w_1 = -2$
- Is this the only solution?

## Examples

---

AND

$x_0$	$x_1$	$x_2$		$t$
1	0	0		0
1	0	1		0
1	1	0		0
1	1	1		1

$$z = w_0x_0 + w_1x_1 + w_2x_2$$

Need:  $w_0 < 0$

Need:  $w_0 + w_2 < 0$

Need:  $w_0 + w_1 < 0$

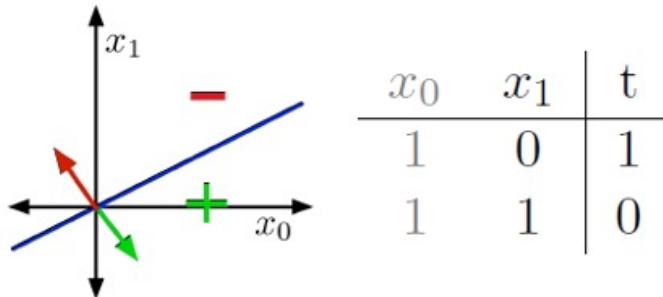
Need:  $w_0 + w_1 + w_2 \geq 0$

Example solution:  $w_0 = -1.5, w_1 = 1, w_2 = 1$

# The Geometric Picture

---

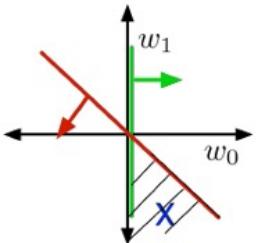
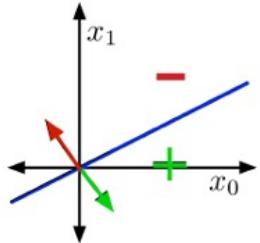
Input Space, or Data Space for NOT example



- Training examples are points
- Weights (hypotheses)  $\mathbf{w}$  can be represented by half-spaces  
 $H_+ = \{\mathbf{x}: \mathbf{w}^\top \mathbf{x} \geq 0\}, H_- = \{\mathbf{x}: \mathbf{w}^\top \mathbf{x} < 0\}$ 
  - The boundaries of these half-spaces pass through the origin (why?)
- The boundary is the decision boundary:  $\{\mathbf{x}: \mathbf{w}^\top \mathbf{x} = 0\}$ 
  - In 2-D, it's a line, but in high dimensions it is a hyperplane
- If the training examples can be perfectly separated by a linear decision rule, we say data is linearly separable.

# The Geometric Picture

## Weight Space



$$w_0 \geq 0$$

$$w_0 + w_1 < 0$$

- Weights (hypotheses)  $\mathbf{w}$  are points
- Each training example  $\mathbf{x}$  specifies a half-space  $\mathbf{w}$  must lie in to be correctly classified:  
 $\mathbf{w}^\top \mathbf{x} \geq 0$  if  $t = 1$ .
- For NOT example:
  - $x_0 = 1, x_1 = 0, t = 1 \iff (w_0, w_1) \in \{\mathbf{w}: w_0 \geq 0\}$
  - $x_0 = 1, x_1 = 1, t = 0 \iff (w_0, w_1) \in \{\mathbf{w}: w_0 + w_1 < 0\}$
- The region satisfying all the constraints is the **feasible region**; if this region is nonempty, the problem is **feasible**, otw it is **infeasible**.

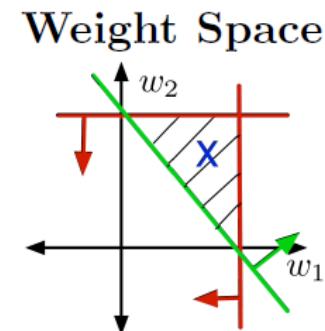
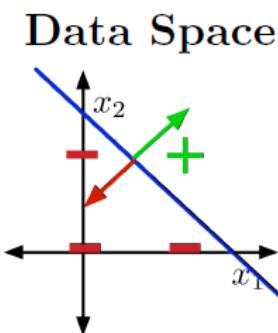
## The Geometric Picture

---

- The **AND** example requires three dimensions, including the dummy one.
- To visualize data space and weight space for a 3-D example, we can look at a 2-D slice.
- The visualizations are similar.
  - Feasible set will always have a corner at the origin.

# The Geometric Picture

## Visualizations of the AND example



- Slice for  $x_0 = 1$  and
- Example sol:  $w_0 = -1.5, w_1 = 1, w_2 = 1$
- Decision boundary:  
 $w_0x_0 + w_1x_1 + w_2x_2 = 0$   
 $\implies -1.5 + x_1 + x_2 = 0$

- Slice for  $w_0 = -1.5$  for the constraints
- $w_0 < 0$
- $w_0 + w_2 < 0$
- $w_0 + w_1 < 0$
- $w_0 + w_1 + w_2 \geq 0$

## Summary | Binary Linear Classifiers

---

- **Summary:** targets  $t \in \{0,1\}$ , inputs  $\mathbf{x} \in \mathbb{R}^{D+1}$  with  $x_0 = 1$ , and model is defined by weights  $\mathbf{w}$  and

$$z = \mathbf{w}^\top \mathbf{x}$$

$$y = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

- How can we find good values for  $\mathbf{w}$ ?
- If training set is linearly separable, we could solve for  $\mathbf{w}$  using linear programming
  - We could also apply an iterative procedure known as the *perceptron algorithm* (but this is primarily of historical interest).
- If it's not linearly separable, the problem is harder
  - Data is almost never linearly separable in real life.

# Towards Logistic Regression

---

Towards Logistic Regression

## Loss Functions

---

- Instead: define loss function then try to minimize the resulting cost function
  - Recall : cost is loss averaged (or summed) over the training set
- Seemingly obvious loss function: **0-1 loss**

$$\begin{aligned}\mathcal{L}_{0-1}(y, t) &= \begin{cases} 0 & \text{if } y = t \\ 1 & \text{if } y \neq t \end{cases} \\ &= \mathbb{I}[y \neq t]\end{aligned}$$

## Attempt 1: 0-1 loss

---

- Usually, the cost  $\mathcal{J}$  is the averaged loss over training examples; for 0-1 loss, this is the **misclassification rate**:

$$\mathcal{J} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}[y^{(i)} \neq t^{(i)}]$$

## Attempt 1: 0-1 loss

---

- Problem: how to optimize? In general, a hard problem (can be NP-hard)
- This is due to the step function (0-1 loss) not being nice (continuous/smooth/convex etc)

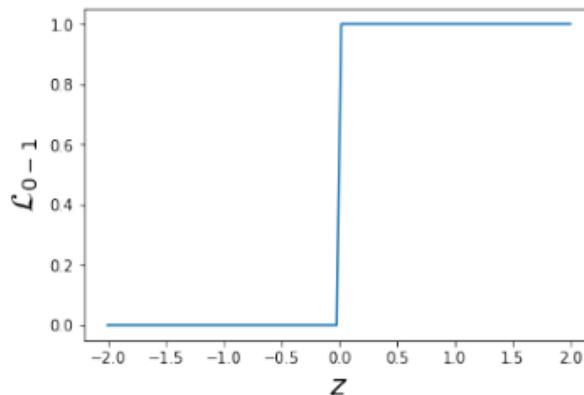
## Attempt 1: 0-1 loss

---

- Minimum of a function will be at its critical points.
- Let's try to find the critical point of 0-1 loss.
- Chain rule:

$$\frac{\partial \mathcal{L}_{0-1}}{\partial w_j} = \frac{\partial \mathcal{L}_{0-1}}{\partial z} \frac{\partial z}{\partial w_j}$$

- But  $\partial \mathcal{L}_{0-1}/\partial z$  is zero everywhere it's defined!



- $\partial \mathcal{L}_{0-1}/\partial w_j = 0$  means that changing the weights by a very small amount probably has no effect on the loss.
- Almost any point has 0 gradient!

## Attempt 2: Linear Regression

---

- Sometimes we can replace the loss function we care about with one which is easier to optimize. This is known as **relaxation** with a smooth **surrogate loss function**.
- One problem with  $\mathcal{L}_{0-1}$ : defined in terms of final prediction, which inherently involves a discontinuity.
- Instead, define loss in terms of  $\mathbf{w}^T \mathbf{x}$  directly.
  - Redo notation for convenience:  $z = \mathbf{w}^T \mathbf{x}$

## Attempt 2: Linear Regression

---

- We already know how to fit a linear regression model. Can we use this instead?

$$z = \mathbf{w}^\top \mathbf{x}$$

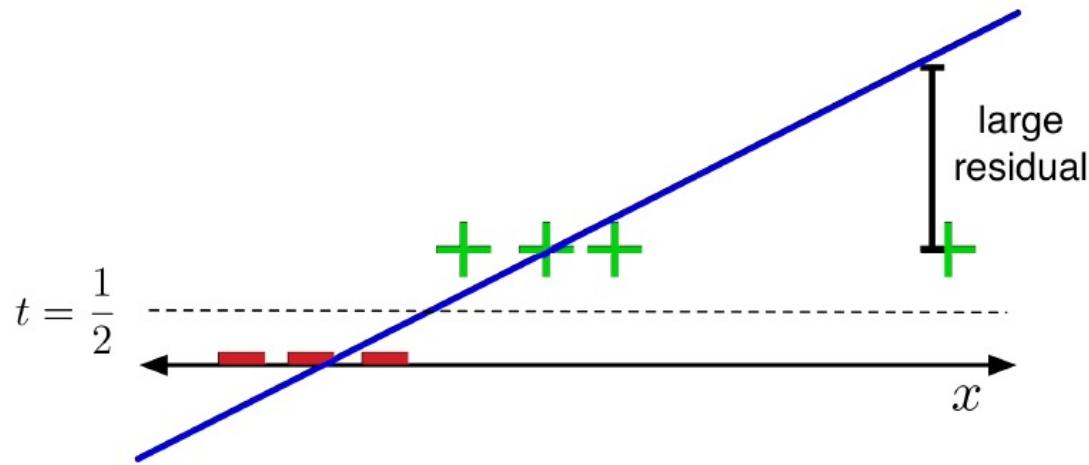
$$\mathcal{L}_{\text{SE}}(z, t) = \frac{1}{2}(z - t)^2$$

- Doesn't matter that the targets are actually binary. Treat them as continuous values.
- For this loss function, it makes sense to make final predictions by thresholding  $z$  at  $\frac{1}{2}$  (why?)

## Attempt 2: Linear Regression

---

**The problem:**



- The loss function hates when you make correct predictions with high confidence!
- If  $t = 1$ , it's more unhappy about  $z = 10$  than  $z = 0$ .

## Attempt 3: Logistic Activation Function

- There's obviously no reason to predict values outside [0, 1]. Let's squash  $y$  into this interval.
- The **logistic function** is a kind of **sigmoid**, or S-shaped function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

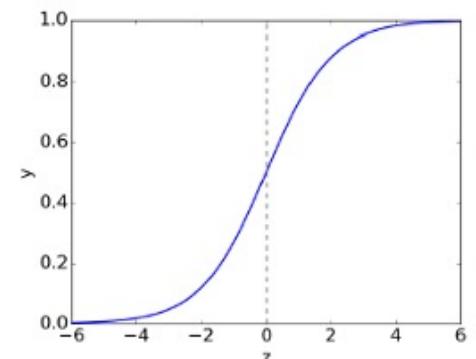
- $\sigma^{-1}(y) = \log(y/(1 - y))$  is called the **logit**.
- A linear model with a logistic nonlinearity is known as **log-linear**:

$$z = \mathbf{w}^\top \mathbf{x}$$

$$y = \sigma(z)$$

$$\mathcal{L}_{\text{SE}}(y, t) = \frac{1}{2} (y - t)^2$$

- Used in this way,  $\sigma$  is called an **activation function**.

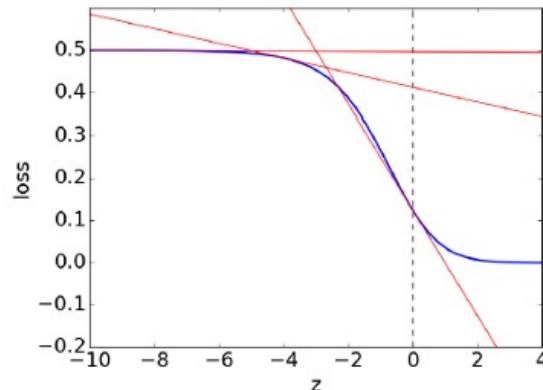


## Attempt 3: Logistic Activation Function

---

**The problem:**

(plot of  $\mathcal{L}_{\text{SE}}$  as a function of  $z$ , assuming  $t = 1$ )



$$\frac{\partial \mathcal{L}}{\partial w_j} = \frac{\partial \mathcal{L}}{\partial z} \frac{\partial z}{\partial w_j}$$

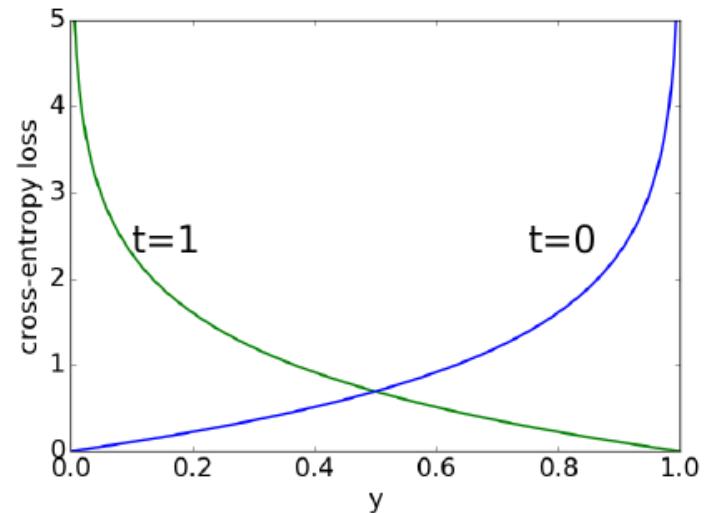
- For  $z \ll 0$ , we have  $\sigma(z) \approx 0$ .
- $\frac{\partial \mathcal{L}}{\partial z} \approx 0$  (check)  $\rightarrow \frac{\partial \mathcal{L}}{\partial w_j} \approx 0 \rightarrow$  derivative w.r.t.  $w_j$  is small  $\rightarrow w_j$  is like a critical point.
- If the prediction is really wrong, you should be far from a critical point (which is your candidate solution).

## Logistic Regression

---

- Because  $y \in [0, 1]$ , we can interpret it as the estimated probability that  $t = 1$ . If  $t = 0$ , then we want to heavily penalize  $y \approx 1$ .
- The pundits who were 99% confident Clinton would win were much more wrong than the ones who were only 90% confident.
- **Cross-entropy loss** (aka log loss) captures this intuition:

$$\begin{aligned}\mathcal{L}_{\text{CE}}(y, t) &= \begin{cases} -\log y & \text{if } t = 1 \\ -\log(1 - y) & \text{if } t = 0 \end{cases} \\ &= -t \log y - (1 - t) \log(1 - y)\end{aligned}$$



# Logistic Regression

---

Logistic regression:

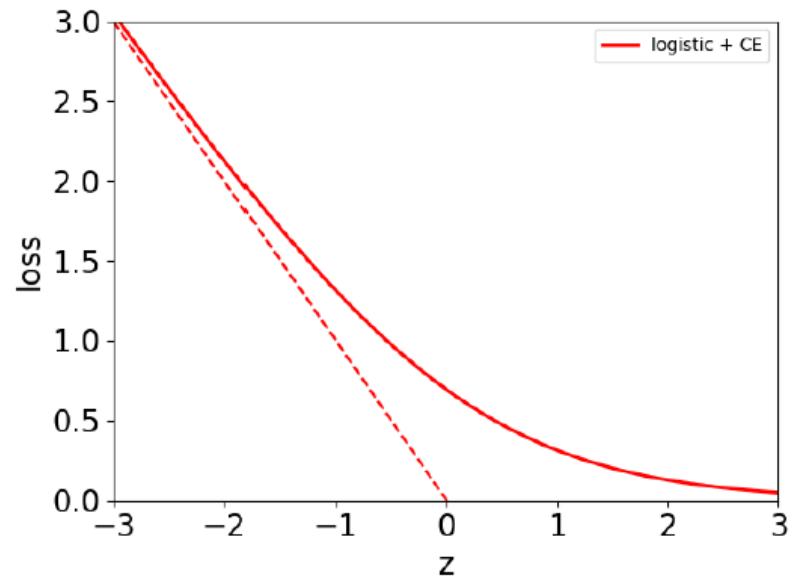
$$z = \mathbf{w}^\top \mathbf{x}$$

$$y = \sigma(z)$$

$$= \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}_{\text{CE}}(y, t) = -t \log y - (1 - t) \log(1 - y)$$

Plot is for target  $t = 1$ .



## Gradient Descent for Logistic Regression

---

- How do we minimize the cost  $\mathcal{J}$  for logistic regression? No direct solution.
  - Taking derivatives of  $\mathcal{J}$  w.r.t.  $\mathbf{w}$  and setting them to 0 doesn't have an explicit solution.
- However, the logistic loss is a **convex function** in  $\mathbf{w}$ , so let's consider the **gradient descent** method from the last lecture.
  - Recall: we **initialize** the weights to something reasonable and repeatedly adjust them in the **direction of steepest descent**.
  - A standard initialization is  $\mathbf{w} = 0$ . (why?)

## Gradient of Logistic Loss

---

Back to logistic regression:

$$\mathcal{L}_{\text{CE}}(y, t) = -t \log y - (1 - t) \log(1 - y)$$

$$y = 1/(1 + e^{-z}) \text{ and } z = \mathbf{w}^\top \mathbf{x}$$

Therefore

$$\begin{aligned}\frac{\partial \mathcal{L}_{\text{CE}}}{\partial w_j} &= \frac{\partial \mathcal{L}_{\text{CE}}}{\partial y} \cdot \frac{\partial y}{\partial z} \cdot \frac{\partial z}{\partial w_j} = \left( -\frac{t}{y} + \frac{1-t}{1-y} \right) \cdot y(1-y) \cdot x_j \\ &= (y - t)x_j\end{aligned}$$

(verify this)

Gradient descent (coordinatewise) update to find the weights of logistic regression:

$$\begin{aligned}w_j &\leftarrow w_j - \alpha \frac{\partial \mathcal{J}}{\partial w_j} \\ &= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)}\end{aligned}$$

# Gradient Descent for Logistic Regression

---

## Comparison of gradient descent updates:

- Linear regression:

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$

- Logistic regression:

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$

- Not a coincidence! These are both examples of **generalized linear models**. But we won't go in further detail.
- Notice  $\frac{1}{N}$  in front of sums due to averaged losses. This is why you need smaller learning rate when cost is summed losses ( $\alpha' = \alpha/N$ ).

## Multiclass Classification

---

Multiclass Classification and Softmax Regression

## Overview

---

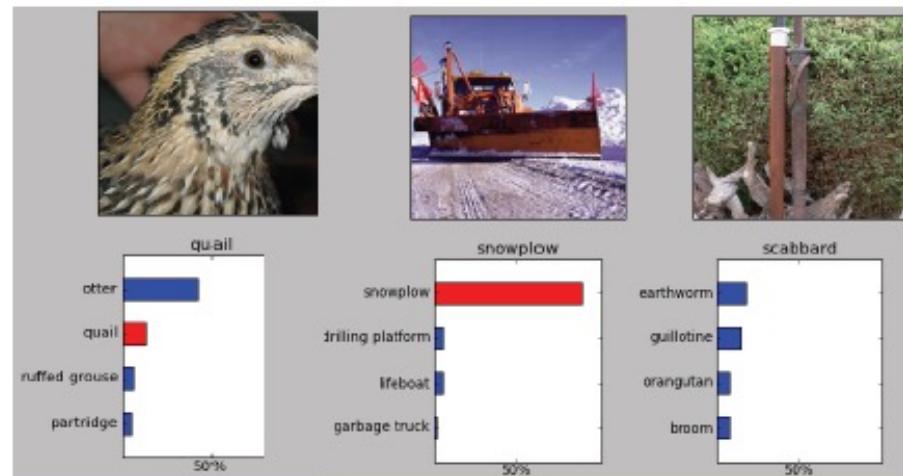
- **Classification**: predicting a discrete-valued target
  - **Binary classification**: predicting a binary-valued target
  - **Multiclass classification**: predicting a discrete-( $>2$ )-valued target
- Examples of multi-class classification
  - Predict the value of a handwritten digit
  - Classify e-mails as spam, travel, work, personal

# Multiclass Classification

---

- Classification tasks with more than two categories:

0 0 0 1 1 1 1 1 2  
2 2 2 2 2 2 2 3 3 3  
3 4 4 4 4 5 5 5 5  
6 6 7 7 7 7 8 8 8  
8 8 8 8 9 9 9 9



## Multiclass Classification

---

- Targets form a discrete set  $\{1, \dots, K\}$ .
- It's often more convenient to represent them as **one-hot vectors**, or a **one-of-K encoding**:

$$\mathbf{t} = (0, \dots, 0, \underbrace{1}_{\text{entry } k \text{ is } 1}, 0 \dots, 0) \in \mathbb{R}^K$$

## Multiclass Linear Classification

---

- We can start with a linear function of the inputs.
- Now there are  $D$  input dimensions and  $K$  output dimensions, so we need  $K \times D$  weights, which we arrange as a **weight matrix  $\mathbf{W}$** .
- Also, we have a  $K$ -dimensional vector  $\mathbf{b}$  of biases.
- A linear function of the inputs:

$$z_k = \sum_{j=1}^D w_{kj}x_j + b_k \quad \text{for } k = 1, 2, \dots, K$$

- We can eliminate the bias  $\mathbf{b}$  by taking  $\mathbf{W} \in \mathbb{R}^{K \times (D+1)}$  and adding a dummy variable  $x_0 = 1$ . So, vectorized:

$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b} \quad \text{or with dummy } x_0 = 1 \quad \mathbf{z} = \mathbf{W}\mathbf{x}$$

## Multiclass Linear Classification

---

- How can we turn this linear prediction into a **one-hot prediction**?
- We can interpret the magnitude of  $z_k$  as a measure of how much the model prefers  $k$  as its prediction.
- If we do this, we should set

$$y_i = \begin{cases} 1 & i = \arg \max_k z_k \\ 0 & \text{otherwise} \end{cases}$$

## Softmax Regression

---

- We need to soften our predictions for the sake of optimization.
- We want soft predictions that are like probabilities, i.e.,  $0 \leq y_k \leq 1$  and  $\sum_k y_k = 1$ .
- A natural activation function to use is the **softmax function**, a multivariable generalization of the logistic function:

$$y_k = \text{softmax}(z_1, \dots, z_K)_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$$

- Outputs can be interpreted as probabilities (positive and sum to 1)
- If  $z_k$  is much larger than the others, then  $\text{softmax}(\mathbf{z})_k \approx 1$  and it behaves like argmax.

## Softmax Regression

---

- If a model outputs a vector of class probabilities, we can use cross-entropy as the loss function:

$$\begin{aligned}\mathcal{L}_{\text{CE}}(\mathbf{y}, \mathbf{t}) &= - \sum_{k=1}^K t_k \log y_k \\ &= -\mathbf{t}^\top (\log \mathbf{y}),\end{aligned}$$

where the log is applied elementwise.

- Just like with logistic regression, we typically combine the softmax and cross-entropy into a **softmax-cross-entropy** function.

## Softmax Regression

---

- Softmax regression (with dummy  $x_0 = 1$ ):

$$\mathbf{z} = \mathbf{W}\mathbf{x}$$

$$\mathbf{y} = \text{softmax}(\mathbf{z})$$

$$\mathcal{L}_{\text{CE}} = -\mathbf{t}^\top (\log \mathbf{y})$$

- Gradient descent updates can be derived for each row of  $\mathbf{W}$ :

$$\frac{\partial \mathcal{L}_{\text{CE}}}{\partial \mathbf{w}_k} = \frac{\partial \mathcal{L}_{\text{CE}}}{\partial z_k} \cdot \frac{\partial z_k}{\partial \mathbf{w}_k} = (y_k - t_k) \cdot \mathbf{x}$$

$$\mathbf{w}_k \leftarrow \mathbf{w}_k - \alpha \frac{1}{N} \sum_{i=1}^N (y_k^{(i)} - t_k^{(i)}) \mathbf{x}^{(i)}$$

- Similar to linear/logistic reg (no coincidence) (verify the update)

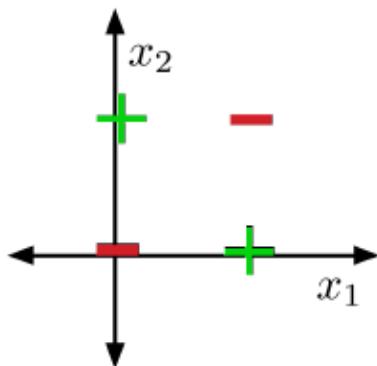
Prove the gradient ?

---

## Limits of Linear Classification

---

Some datasets are not linearly separable, e.g. **XOR**



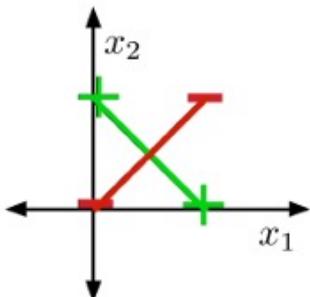
Visually obvious, but how to show this?

## Limits of Linear Classification

---

### Showing that XOR is not linearly separable (proof by contradiction)

- If two points lie in a half-space, line segment connecting them also lie in the same half-space.
- Suppose there were some feasible weights (hypothesis). If the positive examples are in the positive half-space, then the green line segment must be as well.
- Similarly, the red line segment must lie within the negative half-space.



- But the intersection can't lie in both half-spaces. Contradiction!

## Limits of Linear Classification

---

- Sometimes we can overcome this limitation using **feature maps**, just like for linear regression. E.g., for **XOR**:

$$\psi(\mathbf{x}) = \begin{pmatrix} x_1 \\ x_2 \\ x_1x_2 \end{pmatrix}$$

$x_1$	$x_2$	$\psi_1(\mathbf{x})$	$\psi_2(\mathbf{x})$	$\psi_3(\mathbf{x})$	$t$
0	0	0	0	0	0
0	1	0	1	0	1
1	0	1	0	0	1
1	1	1	1	1	0

- This is linearly separable. (Try it!)

## Next time...

---

Feature maps are hard to design well, so next time we'll see how to *learn* nonlinear feature maps directly using neural networks...

