



**Xi'an Jiaotong-Liverpool University**

**西交利物浦大学**

**INT305 Machine Learning**

**Lecture 7**

**Decision Trees & Bias-Variance Decomposition**

**Siyue Yu**

**Department Intelligence Science**

**[siyue.yu02@xjtu.edu.cn](mailto:siyue.yu02@xjtu.edu.cn)**

# Today


---

- **Decision Trees**
  - Simple but powerful learning algorithm
  - Used widely in Kaggle competitions
  - Let us motivate concepts from information theory (entropy, mutual information, etc.)
- **Bias-variance decomposition**
  - Let us motivate methods for combining different classifiers.

# Decision Trees

---


- Make predictions by splitting on features according to a tree structure.



# Decision Trees



---

- Make predictions by splitting on features according to a tree structure.




# Decision Trees— Continuous Features

- Split *continuous features* by checking whether that feature is greater than or less than some threshold.
- Decision boundary is made up of axis-aligned planes.



# Decision Trees


---



- Internal nodes test a feature
- Branching is determined by the feature value
- Leaf nodes are outputs (predictions)


# Decision Trees—Classification and Regression

- Each path from root to a leaf defines a region  $R_m$  of input space
- Let  $\{(x^{(m_1)}, t^{(m_1)}), \dots, (x^{(m_k)}, t^{(m_k)})\}$  be the training examples that fall into  $R_m$
- Classification tree (we will focus on this):
  - discrete output
  - leaf value  $y^m$  typically set to the most common value in  $\{t^{(m_1)}, \dots, t^{(m_k)}\}$
- Regression tree:
  - Continuous output
  - leaf value  $y^m$  typically set to the mean value in  $\{t^{(m_1)}, \dots, t^{(m_k)}\}$



# Decision Trees—Discrete Features

- Will I wait at this restaurant for dinner?



# Decision Trees—Discrete Features

- Split discrete features into a partition of possible values.

Example	Input Attributes										Goal <i>WillWait</i>
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	
$x_1$	Yes	No	No	Yes	Some	\$\$\$	No	Yes	French	0-10	$y_1 = \text{Yes}$
$x_2$	Yes	No	No	Yes	Full	\$	No	No	Thai	30-60	$y_2 = \text{No}$
$x_3$	No	Yes	No	No	Some	\$	No	No	Burger	0-10	$y_3 = \text{Yes}$
$x_4$	Yes	No	Yes	Yes	Full	\$	Yes	No	Thai	10-30	$y_4 = \text{Yes}$
$x_5$	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	$y_5 = \text{No}$
$x_6$	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0-10	$y_6 = \text{Yes}$
$x_7$	No	Yes	No	No	None	\$	Yes	No	Burger	0-10	$y_7 = \text{No}$
$x_8$	No	No	No	Yes	Some	\$\$	Yes	Yes	Thai	0-10	$y_8 = \text{Yes}$
$x_9$	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	$y_9 = \text{No}$
$x_{10}$	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10-30	$y_{10} = \text{No}$
$x_{11}$	No	No	No	No	None	\$	No	No	Thai	0-10	$y_{11} = \text{No}$
$x_{12}$	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30-60	$y_{12} = \text{Yes}$


1.	Alternate: whether there is a suitable alternative restaurant nearby.
2.	Bar: whether the restaurant has a comfortable bar area to wait in.
3.	Fri/Sat: true on Fridays and Saturdays.
4.	Hungry: whether we are hungry.
5.	Patrons: how many people are in the restaurant (values are None, Some, and Full).
6.	Price: the restaurant's price range (\$, \$\$, \$\$\$).
7.	Raining: whether it is raining outside.
8.	Reservation: whether we made a reservation.
9.	Type: the kind of restaurant (French, Italian, Thai or Burger).
10.	WaitEstimate: the wait estimated by the host (0-10 minutes, 10-30, 30-60, >60).

Features:

## Attempts


---

The drawing below shows a dataset. Each example in the dataset has two inputs features  $x$  and  $y$ , and maybe classified as a positive example (labelled +) or a negative example (labelled -). Draw a decision tree which correctly classifies each example in the dataset.



# Attempts

---



# Learning Decision Trees

---

- For any training set we can construct a decision tree that has exactly one leaf for every training point, but it probably won't generalize.
  - Decision trees are universal function approximators.
- But, finding the smallest decision tree that correctly classifies a training set is NP complete.
  - If you are interested, check: Hyafil & Rivest'76.
- So, how do we construct a useful decision tree?

# Learning Decision Trees


---

- Resort to a **greedy heuristic**:
  - Start with the whole training set and an empty decision tree.
  - Pick a feature and candidate split that would most reduce the loss.
  - Split on that feature and recurse on subpartitions.
- Which loss should we use?
  - Let's see if misclassification rate is a good loss.

# Choosing a Good Split

---


- Consider the following data. Let's split on width.



# Choosing a Good Split

---

- Recall: classify by majority.




- A and B have the same misclassification rate, so which is the best split? Vote!

# Choosing a Good Split

---

- A feels like a better split, because the left-hand region is very certain about whether the fruit is an orange.



- Can we quantify this?

# Choosing a Good Split

---

- How can we quantify uncertainty in prediction for a given leaf node?
  - If all examples in leaf have same class: good, low uncertainty
  - If each class has same amount of examples in leaf: bad, high uncertainty
- **Idea:** use counts at leaves to define probability distributions; use a probabilistic notion of uncertainty to decide splits.
- A brief detour through information theory...

# Quantifying Uncertainty

---

- The **entropy** of a discrete random variable is a number that quantifies the **uncertainty** inherent in its possible outcomes.
- The mathematical definition of entropy that we give in a few slides may seem arbitrary, but it can be motivated axiomatically.
  - If you're interested, check: *Information Theory* by Robert Ash.
- To explain entropy, consider flipping two different coins...

# We Flip Two Different Coins



---

Sequence 1:

0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 ... ?

Sequence 2:



0 1 0 1 0 1 1 1 0 1 0 0 1 1 1 0 1 0 1 ... ?



# Quantifying Uncertainty

- The entropy of a loaded coin with probability  $p$  of heads is given by

$$-p\log_2(p) - (1-p)\log_2(1-p)$$



$$-\frac{8}{9} \log_2 \frac{8}{9} - \frac{1}{9} \log_2 \frac{1}{9} \approx \frac{1}{2}$$


$$-\frac{4}{9} \log_2 \frac{4}{9} - \frac{5}{9} \log_2 \frac{5}{9} \approx 0.99$$

- Notice: the coin whose outcomes are more certain has a lower entropy
- In the extreme case  $p = 0$  or  $p = 1$ , we were certain of the outcome before observing. So, we gained no certainty by observing it, i.e., entropy is 0.

# Quantifying Uncertainty

---

- Can also think of **entropy** as the expected information content of a random draw from a probability distribution.



- Claude Shannon showed: you cannot store the outcome of a random draw using fewer expected bits than the entropy without losing information.
- So units of entropy are **bits**; a fair coin flip has 1 bit of entropy.

# Entropy

---

- More generally, the entropy of a discrete random variable  $Y$  is given by

$$H(Y) = - \sum_{y \in Y} p(y) \log_2 p(y)$$

- “High Entropy”
  - Variable has a uniform like distribution over many outcomes
  - Flat histogram
  - Values sampled from it are less predictable
- “Low Entropy”
  - Distribution is concentrated on only a few outcomes
  - Histogram is concentrated in a few areas
  - Values sampled from it are more predictable

# Entropy

---

- Suppose we observe partial information  $X$  about a random variable  $Y$ 
  - For example,  $X = \text{sign}(Y)$ .
- We want to work towards a definition of the expected amount of information that will be conveyed about  $Y$  by observing  $X$ .
  - Or equivalently, the expected reduction in our uncertainty about  $Y$  after observing  $X$ .

# Entropy of a Joint Distribution

- Example:  $X = \{\text{Raining, Not raining}\}, Y = \{\text{Cloudy, Not cloudy}\}$

		Cloudy	Not Cloudy
Raining	24/100	1/100	
Not Raining	25/100	50/100	

$$\begin{aligned} H(X, Y) &= - \sum_{x \in X} \sum_{y \in Y} p(x, y) \log_2 p(x, y) \\ &= -\frac{24}{100} \log_2 \frac{24}{100} - \frac{1}{100} \log_2 \frac{1}{100} - \frac{25}{100} \log_2 \frac{25}{100} - \frac{50}{100} \log_2 \frac{50}{100} \\ &\approx 1.56 \text{ bits} \end{aligned}$$

# Specific Conditional Entropy

---

- Example:  $X = \{\text{Raining, Not raining}\}$ ,  $Y = \{\text{Cloudy, Not cloudy}\}$

	Cloudy	Not Cloudy
Raining	24/100	1/100
Not Raining	25/100	50/100

- What is the entropy of cloudiness  $Y$ , given that it is raining?

$$\begin{aligned} H(Y|X = x) &= - \sum_{y \in Y} p(y|x) \log_2 p(y|x) \\ &= -\frac{24}{25} \log_2 \frac{24}{25} - \frac{1}{25} \log_2 \frac{1}{25} \\ &\approx 0.24 \text{ bits} \end{aligned}$$

- We used:  $p(y|x) = \frac{p(x,y)}{p(x)}$  and  $p(x) = \sum_y p(x,y)$  (sum in a row)

# Conditional Entropy

---

	Cloudy	Not Cloudy
Raining	24/100	1/100
Not Raining	25/100	50/100

- The expected conditional entropy:

$$\begin{aligned} H(Y|X) &= \sum_{x \in X} p(x)H(Y|X=x) \\ &= -\sum_{x \in X} \sum_{y \in Y} p(x,y)\log_2 p(y|x) \end{aligned}$$

# Conditional Entropy

---

- Example:  $X = \{\text{Raining, Not raining}\}$ ,  $Y = \{\text{Cloudy, Not cloudy}\}$

		Cloudy	Not Cloudy
Raining	24/100	1/100	
Not Raining	25/100	50/100	

- What is the entropy of cloudiness  $Y$ , **given the knowledge of whether or not it is raining?**

$$\begin{aligned} H(Y|X) &= \sum_{x \in X} p(x)H(Y|X=x) \\ &= \frac{1}{4}H(\text{cloudiness|is raining}) + \frac{3}{4}H(\text{cloudiness|not rainig}) \\ &\approx 0.75\text{bits} \end{aligned}$$

# Conditional Entropy

---

- Some useful properties
  - $H$  is always non-negative
  - Chain rule:  $H(X, Y) = H(X|Y) + H(Y) = H(Y|X) + H(X)$
  - If  $X$  and  $Y$  independent, then  $X$  does not affect our uncertainty about  $Y$ :  $H(Y|X) = H(Y)$
  - But knowing  $Y$  makes our knowledge of  $Y$  certain:  $H(Y|Y) = 0$
  - By knowing  $X$ , we can only decrease uncertainty about  $Y$ :  $H(Y|X) \leq H(Y)$

# Information Gain

---

	Cloudy	Not Cloudy
Raining	24/100	1/100
Not Raining	25/100	50/100

- How much more certain am I about whether it's cloudy if I'm told whether it is raining? My uncertainty in  $Y$  minus my expected uncertainty that would remain in  $Y$  after seeing  $X$ .
- This is called the **information gain**  $IG(Y|X)$  in  $Y$  due to  $X$ , or the **mutual information** of  $Y$  and  $X$

$$IG(Y|X) = H(Y) - H(Y|X) \tag{1}$$

- If  $X$  is completely uninformative about  $Y$ :  $IG(Y|X) = 0$
- If  $X$  is completely informative about  $Y$ :  $IG(Y|X) = H(Y)$

## Revisiting Our Original Example


---

- Information gain measures the informativeness of a variable, which is exactly what we desire in a decision tree split!
- The information gain of a split: how much information (over the training set) about the class label  $Y$  is gained by knowing which side of a split you're on.

# Revisiting Our Original Example

---


- What is the information gain of split B? Not terribly informative...



- Root entropy of class outcome:  $H(Y) = -\frac{2}{7} \log_2(\frac{2}{7}) - \frac{5}{7} \log_2(\frac{5}{7}) \approx 0.86$
- Leaf conditional entropy of class outcome:  $H(Y|left) \approx 0.81, H(Y|right) \approx 0.92$
- $IG(split) \approx 0.86 - (\frac{4}{7} \cdot 0.81 + \frac{3}{7} \cdot 0.92) \approx 0.006$



# Revisiting Our Original Example

- What is the information gain of split A? Very informative!



- Root entropy of class outcome:  $H(Y) = -\frac{2}{7} \log_2(\frac{2}{7}) - \frac{5}{7} \log_2(\frac{5}{7}) \approx 0.86$
- Leaf conditional entropy of class outcome:  $H(Y|left) \approx 0, H(Y|right) \approx 0.97$
- $IG(split) \approx 0.86 - (\frac{2}{7} \cdot 0 + \frac{5}{7} \cdot 0.97) \approx 0.17!$

# Constructing Decision Trees



- At each level, one must choose:
  1. Which feature to split.
  2. Possibly where to split it.
- Choose them based on how much information we would gain from the decision! (choose feature that gives the highest gain)

# Decision Tree Construction Algorithm

---

- Simple, greedy, recursive approach, builds up tree node-by-node
  1. Pick a feature to split at a non-terminal node
  2. Split examples into groups based on feature value
  3. For each group:
    - If no examples – return majority from parent
    - Else if all examples in same class – return class
    - Else loop to step 1
- Terminates when all leaves contain only examples in the same class or are empty.

# Back to Our Example


Example	Input Attributes										Goal <i>WillWait</i>
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	
$x_1$	Yes	No	No	Yes	Some	\$\$\$	No	Yes	French	0–10	$y_1 = \text{Yes}$
$x_2$	Yes	No	No	Yes	Full	\$	No	No	Thai	30–60	$y_2 = \text{No}$
$x_3$	No	Yes	No	No	Some	\$	No	No	Burger	0–10	$y_3 = \text{Yes}$
$x_4$	Yes	No	Yes	Yes	Full	\$	Yes	No	Thai	10–30	$y_4 = \text{Yes}$
$x_5$	Yes	No	Yes	No	Full	\$\$\$	No	Yes	French	>60	$y_5 = \text{No}$
$x_6$	No	Yes	No	Yes	Some	\$\$	Yes	Yes	Italian	0–10	$y_6 = \text{Yes}$
$x_7$	No	Yes	No	No	None	\$	Yes	No	Burger	0–10	$y_7 = \text{No}$
$x_8$	No	No	No	Yes	Some	\$\$	Yes	Yes	Thai	0–10	$y_8 = \text{Yes}$
$x_9$	No	Yes	Yes	No	Full	\$	Yes	No	Burger	>60	$y_9 = \text{No}$
$x_{10}$	Yes	Yes	Yes	Yes	Full	\$\$\$	No	Yes	Italian	10–30	$y_{10} = \text{No}$
$x_{11}$	No	No	No	No	None	\$	No	No	Thai	0–10	$y_{11} = \text{No}$
$x_{12}$	Yes	Yes	Yes	Yes	Full	\$	No	No	Burger	30–60	$y_{12} = \text{Yes}$

- |     |   |
|-----|---|
| 1.  | Alternate: whether there is a suitable alternative restaurant nearby.             |
| 2.  | Bar: whether the restaurant has a comfortable bar area to wait in.                |
| 3.  | Fri/Sat: true on Fridays and Saturdays.   |
| 4.  | Hungry: whether we are hungry.  |
| 5.  | Patrons: how many people are in the restaurant (values are None, Some, and Full). |
| 6.  | Price: the restaurant's price range (\$, \$\$, \$\$\$).                           |
| 7.  | Raining: whether it is raining outside.   |
| 8.  | Reservation: whether we made a reservation.                                       |
| 9.  | Type: the kind of restaurant (French, Italian, Thai or Burger).                   |
| 10. | WaitEstimate: the wait estimated by the host (0-10 minutes, 10-30, 30-60, >60).   |

Features:

[from: Russell & Norvig]

# Feature Selection



$$IG(Y) = H(Y) - H(Y|X)$$

$$IG(type) = 1 - \left[ \frac{2}{12} H(Y|\text{Fr.}) + \frac{2}{12} H(Y|\text{It.}) + \frac{4}{12} H(Y|\text{Thai}) + \frac{4}{12} H(Y|\text{Bur.}) \right] = 0$$

$$IG(Patrons) = 1 - \left[ \frac{2}{12} H(0,1) + \frac{4}{12} H(1,0) + \frac{6}{12} H\left(\frac{2}{6}, \frac{4}{6}\right) \right] \approx 0.541$$

# What Makes a Good Tree?

---

- Not too small: need to handle important but possibly subtle distinctions in data
- Not too big:
  - Computational efficiency (avoid redundant, spurious attributes)
  - Avoid over-fitting training examples
  - Human interpretability
- “[Occam’s Razor](#)”: find the simplest hypothesis that fits the observations
  - Useful principle, but hard to formalize (how to define simplicity?)
  - See Domingos, 1999, “The role of Occam’s razor in knowledge discovery”
- We desire small trees with informative nodes near the root

# Decision Tree Miscellany

---

- Problems:
  - You have exponentially less data at lower levels
  - Too big of a tree can overfit the data
  - Greedy algorithms don't necessarily yield the global optimum
- Handling continuous attributes
  - Split based on a threshold, chosen to maximize information gain

# Comparison to some other classifiers

---

Advantages of decision trees over KNNs and neural nets

- Simple to deal with discrete features, missing values, and poorly scaled data
- Fast at test time
- More interpretable

Advantages of KNNs over decision trees

- Few hyperparameters
- Can incorporate interesting distance measures (e.g. shape contexts)

Advantages of neural nets over decision trees

- Able to handle attributes/features that interact in very complex ways (e.g. pixels)

# Comparison to some other classifiers




---

- We've seen many classification algorithms.
- We can combine multiple classifiers into an **ensemble**, which is a set of predictors whose individual decisions are combined in some way to classify new examples
  - E.g., (possibly weighted) majority vote
- For this to be nontrivial, the classifiers must differ somehow, e.g.
  - Different algorithm
  - Different choice of hyperparameters
  - Trained on different data
  - Trained with different weighting of the training examples
- Next lecture, we will study some specific ensembling techniques.

- 
- Today, we deepen our understanding of generalization through a bias-variance decomposition.
    - This will help us understand ensembling methods.

# Bias-Variance Decomposition


- Recall that overly simple models underfit the data, and overly complex models overfit.



- We can quantify this effect in terms of the bias/variance decomposition.
  - Bias and variance of what?


# Bias-Variance Decomposition: Basic Setup

- Suppose the training set  $\mathcal{D}$  consists of pairs  $(x_i, t_i)$  sampled **independent and identically distributed (i.i.d)** from a single **data generating distribution**  $p_{\text{sample}}$ .
- Pick a fixed query point  $\mathbf{x}$  (denoted with a green x).
- Consider an experiment where we sample lots of training sets independently from  $p_{\text{sample}}$ .




# Bias-Variance Decomposition: Basic Setup


- Let's run our learning algorithm on each training set, and compute its prediction  $y$  at the query point  $\mathbf{x}$ .
- we can view  $y$  as a random variable, where the randomness comes from the choice of training set.
- The classification accuracy is determined by the distribution of  $y$ .



$$y = \bullet$$




$$y = \bullet$$



$$y = \bullet$$

# Bias-Variance Decomposition: Basic Setup

Here is the analogous setup for regression:



Since  $y$  is a random variable, we can talk about its expectation, variance, etc.

# Bias-Variance Decomposition: Basic Setup

---

- Recap of basic setup:
  - Fix a query point  $\mathbf{x}$ .
  - Repeat:
    - Sample a random training dataset  $\mathcal{D}$  i.i.d. from the data generating distribution  $p_{\text{sample}}$ .
    - Run the learning algorithm on  $\mathcal{D}$  to get a prediction  $y$  at  $\mathbf{x}$ .
    - Sample the (true) target from the conditional distribution  $p(t|\mathbf{x})$ .
    - Compute the loss  $L(y, t)$ .
- Notice:  $y$  is independent of  $t$ .
- This gives a distribution over the loss at  $\mathbf{x}$ , with expectation  $\mathbb{E}[L(y, t)|\mathbf{x}]$ .
- For each query point  $\mathbf{x}$ , the expected loss is different. We are interested in minimizing the expectation of this with respect to  $\mathbf{x} \sim p_{\text{sample}}$ .

# Bayes Optimality

---

- For now, focus on squared error loss,  $L(y, t) = \frac{1}{2}(y - t)^2$ .
- A first step: suppose we knew the conditional distribution  $p(t|\mathbf{x})$ . What value  $y$  should we predict?
  - Here, we are treating  $t$  as a random variable and choosing  $y$ .
- **Claim:**  $y_* = \mathbb{E}[t|\mathbf{x}]$  is the best possible prediction.
- **Proof:**

$$\begin{aligned}\mathbb{E}[(y - t)^2 | \mathbf{x}] &= \mathbb{E}[y^2 - 2yt + t^2 | \mathbf{x}] \\ &= y^2 - 2y\mathbb{E}[t|\mathbf{x}] + \mathbb{E}[t^2|\mathbf{x}] \\ &= y^2 - 2y\mathbb{E}[t|\mathbf{x}] + \mathbb{E}[t|\mathbf{x}]^2 + \text{Var}[t|\mathbf{x}] \\ &= y^2 - 2yy_* + y_*^2 + \text{Var}[t|\mathbf{x}] \\ &= (y - y_*)^2 + \text{Var}[t|\mathbf{x}]\end{aligned}$$

# Bayes Optimality

---

$$\mathbb{E}[(y - t)^2 | \mathbf{x}] = (y - y_*)^2 + \text{Var}[t | \mathbf{x}]$$

- The first term is nonnegative, and can be made 0 by setting  $y = y_*$ .
- The second term corresponds to the inherent unpredictability, or **noise**, of the targets, and is called the **Bayes error**.
  - This is the best we can ever hope to do with any learning algorithm. An algorithm that achieves it is **Bayes optimal**.
  - Notice that this term doesn't depend on  $y$ .
- This process of choosing a single value  $y_*$  based on  $p(t|\mathbf{x})$  is an example of **decision theory**.

# Bayes Optimality

---

- Now return to treating  $y$  as a random variable (where the randomness comes from the choice of dataset).
- We can decompose out the expected loss (suppressing the conditioning on  $\mathbf{x}$  for clarity):

$$\begin{aligned}\mathbb{E}[(y - t)^2] &= \mathbb{E}[(y - y_\star)^2] + \text{Var}(t) \\ &= \mathbb{E}[y_\star^2 - 2y_\star y + y^2] + \text{Var}(t) \\ &= y_\star^2 - 2y_\star \mathbb{E}[y] + \mathbb{E}[y^2] + \text{Var}(t) \\ &= y_\star^2 - 2y_\star \mathbb{E}[y] + \mathbb{E}[y]^2 + \text{Var}(y) + \text{Var}(t) \\ &= \underbrace{(y_\star - \mathbb{E}[y])^2}_{\text{bias}} + \underbrace{\text{Var}(y)}_{\text{variance}} + \underbrace{\text{Var}(t)}_{\text{Bayes error}}\end{aligned}$$

# Bayes Optimality

---


$$\mathbb{E}[(y - t)^2] = \underbrace{(y_\star - \mathbb{E}[y])^2}_{\text{bias}} + \underbrace{\text{Var}(y)}_{\text{variance}} + \underbrace{\text{Var}(t)}_{\text{Bayes error}}$$

- We just split the expected loss into three terms:
  - **bias**: how wrong the expected prediction is (corresponds to underfitting)
  - **variance**: the amount of variability in the predictions (corresponds to overfitting)
  - Bayes error: the inherent unpredictability of the targets
- Even though this analysis only applies to squared error, we often loosely use “bias” and “variance” as synonyms for “underfitting” and “overfitting”.

# Bias and Variance

---

- Throwing darts = predictions for each draw of a dataset



- Be careful, what doesn't this capture?
  - We average over points  $x$  from the data distribution.